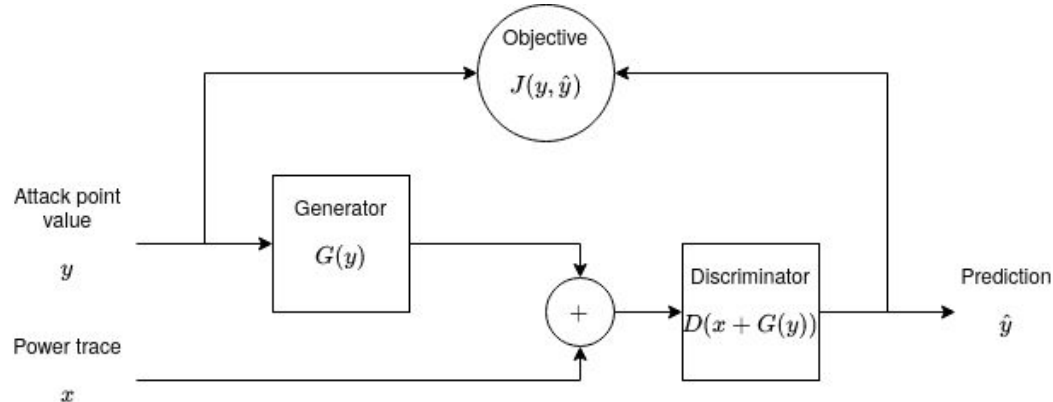


Experiment overview

- 1) Compile system with generator and discriminator
 - a) Variety of generator architectures
 - b) Pretrained discriminator from Google codebase
- 2) Train discriminator to minimize objective J
- 3) Train generator to maximize J
- 4) Again train discriminator to minimize J



Takeaways

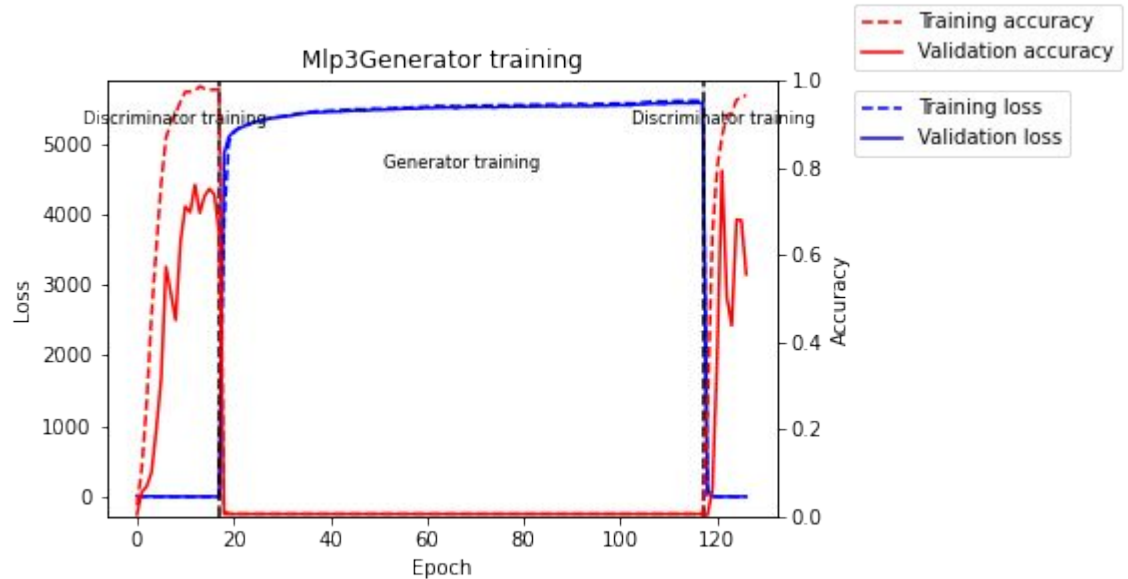
- Easy to disrupt performance of static discriminator
 - Constant offset to trace is sufficient
 - Tends to collapse towards always predicting 1 or few keys
- Re-trained discriminator performance is higher than without discriminator
 - Generator 'gives away' the key
 - Need to re-think generator input, objective
- Basic key -> trace generator has lots of parameters, not very general
 - Could consider recurrent generator, generator controlling sinusoid characteristics
- Currently unclear which generator is best
 - Need to change setup until generator is no longer giving away key

Next steps

- Consider different objective functions
 - Minimize all discriminator outputs, not just for correct key
 - Minimize discriminator gradient magnitude
- Consider different generator inputs
 - Random input instead of correct attack point
 - In practice, only 1 key per generator – could try inputting plaintext instead.
 - Likely solves problem of giving away key
 - Need to find a different dataset
- Try to find better generator architectures
 - Convolutional architectures typical for GANs
 - Sum of sinusoids
 - Recurrent

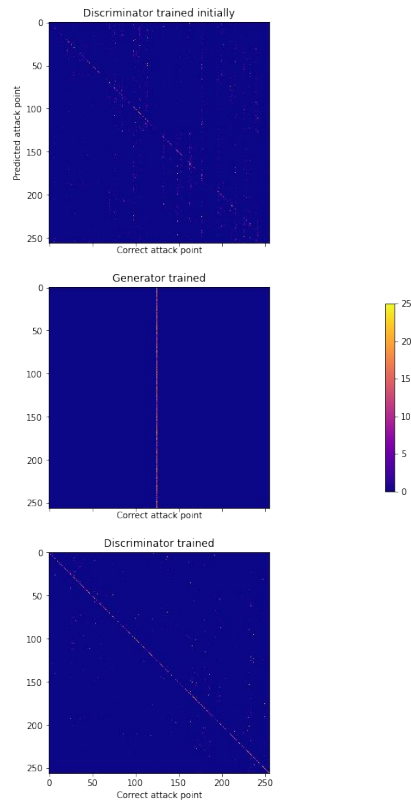
Measurements taken (3-layer MLP generator)

- Loss/accuracy during training



Measurements taken (3-layer MLP generator)

- Confusion matrix after 3 training phases



Measurements taken (3-layer MLP generator)

- Traces before/after application of generator

